

# DICTIONARY DATA STRUCTURES - HASHING

## Open Addressing

- Analysis
- Probing
- Unsuccessful Find
- Successful Find

## Bloom Filters

- Motivation
- Implementation
- Analysis
- Applications.

# TERMINOLOGY

- The technique of chaining elements that hash into the same slot is referred to by different names:
  - Separate Chaining
    - for obvious reasons
  - Open Hashing
    - because number of elements is not limited by table size
  - Closed Address Hashing
    - because the location of the bucket (i.e. the address) of an element is fixed
      - Consider an element  $e$  that is added, removed, and added again:
        - it will get added to the same bucket.

# OPEN ADDRESSING (A.K.A. CLOSED HASHING)

- Fixed Space
  - Fixed Table size and
  - each table location can contain only one element
- Addressing by Hashing
  - Same as in Separate Chaining
- Probing (for a vacant location) in case of collision

## OPEN ADDRESSING (A.K.A. CLOSED HASHING)

○ `add(Element e, Hashtable T)`    // Generic procedure

// e.key is key; h is hash function

`a = h(e.k);`

`if T[a] is empty then { T[a]=e; return; }`

`j=0;`

`repeat {`

`j++;`

`b = getNextAddr(a,k,j);`

`} until (T[b] is empty)`    // Will this terminate?

`T[b] = e;`

# OPEN ADDRESSING – PROBING SCHEMES

// m denotes table size; typically m is chosen to be prime

- Linear Probing:

$\text{getNextAddr}(a, k, j) \{ \text{return } (a+j) \bmod m; \}$

- Quadratic Probing

$\text{getNextAddr}(a, k, j) \{ \text{return } (a+j^2) \bmod m; \}$

- Exponential Probing

$\text{getNextAddr}(a, k, j) \{ \text{return } (a+2^j) \bmod m; \}$

- Double Hashing

$\text{getNextAddr}(a, k, j) \{ \text{return } a+j*h_2(k) \bmod m; \}$

//  $h_2(k)$  is the secondary hash function

//  $h_2(k)$  must be non-zero

// e.g.  $h_2(k) = q - (k \bmod q)$  for some prime  $q < m$

# OPEN ADDRESSING

- Implementation Caveat:
  - add as defined may not terminate!
    - Must check whether all *m* locations have been probed
      - Could be expensive!
    - Alternatively, may use a count of non-empty locations.
      - Will work only if the probing sequence covers all locations

**Exercise:** Handle termination: use simple heuristics(s). **End of Exercise.**

# OPEN ADDRESSING

- Define find.
  - Similar to add:
    - hash
    - if element found return it;
    - if empty return INVALID;
    - otherwise probe until element found or empty slot.
      - return accordingly.
  - Termination?

# OPEN ADDRESSING

- How is deletion done?
  - Deleted slots must be marked *deleted*
    - **deleted** flag different from **empty** flag for probing procedure to work
      - **find** will treat **deleted** slots as **empty** slots
  - This won't allow re-use of **deleted** slots
    - How do you recover deleted slots?
      - **add** can be modified to fill in any *deleted* slot encountered in a probing sequence
        - This may not cover all deleted slots
      - **delete** can be implemented such that subsequent entries in a probing sequence are pulled in.
- How does your deletion scheme affect further probes?



# OPEN ADDRESSING – ANALYSIS OF PROBING

- Probing sequence:
  - Sequence of slots generated:  $S[k,0], S[k,1], \dots S[k,m]$
- Probing requirements:
  - Utilization:
    - The probing sequence must be a permutation of  $0, 1, \dots, m-1$
  - Uniform hashing assumption:
    - Requires that each key may result in any of the  $m!$  probing sequences

## OPEN ADDRESSING – ANALYSIS OF PROBING [2]

### ○ Linear Probing

- Slot for the  $j^{\text{th}}$  probe in a table of size  $m$   
 $S[k,j] = (h(k) + j) \bmod m$
- Long runs of occupied slots build up
  - If an empty slot is preceded by  $j$  full slots,
  - then the probability this slot is the next one filled is  $(j+1)/m$
  - instead of  $1/m$  (in a table of size  $m$ )
- Effect known as (Primary) clustering
- This is not a good approximation of uniform hashing

# OPEN ADDRESSING – ANALYSIS OF PROBING [3]

## ○ Quadratic Probing

- Slot for the  $j^{\text{th}}$  probe in a table of size  $m$   
 $S[k,j] = (h(k) + j^2) \bmod m$
- Clustering effect milder than Linear probing
  - Effect known as secondary clustering
- But the sequence of slots probed is still dependent on the initial slot (decided by the key)
  - i.e only  $m$  distinct sequences are explored

## ○ Generalize:

- $S[k,j] = (h(k) + a*j + b*j^2) \bmod m$

Exercise: Can you choose  $a$ ,  $b$ , and  $m$  such that all slots are utilized?

Exercise: Repeat (very similar) analysis for Exponential Probing.

# OPEN ADDRESSING – ANALYSIS OF PROBING [4]

## ○ Double Hashing

- Slot for the  $j^{\text{th}}$  probe in a table of size  $m$   
 $S[k,j] = (h_1(k) + j \cdot h_2(k)) \bmod m$
- Probing sequence depends on  $k$  in two ways
  - So, probing sequence depends not only on initial slot  
i.e.  $m \cdot m$  probing sequences can be used.

This results in behavior closer to uniform hashing

- If  $\gcd(h_2(k), m) = d$  for some key  $k$ ,
  - then the sequence will explore only  $(1/d) \cdot m$  slots
  - Why?
  - So, choose (for instance):
    - $m$  as a prime, and ensure  $h_2(k)$  is always  $< m$
- Can you extend this to a sequence of hashes  $h_1(k), h_2(k), h_3(k), \dots$  ?

## OPEN ADDRESSING – ANALYSIS - UNSUCCESSFUL FIND

- Given: open-address table with load factor  $\alpha = n/m < 1$
- Assumption: Uniform Hashing
- Expected number of probes in an unsuccessful find is at most  $1/(1 - \alpha)$
- Proof:
  - Last probed slot is empty; all previous probed slots are non-empty but do not contain the given key
  - Define  $p_j$  as the probability that exactly  $j$  probes access non-empty slots
    - Then the expected number of probes is  $1 + \sum_{j=1}^{\infty} j \cdot p_j$
  - If  $q_j$  is defined as the probability that at least  $j$  probes access non-empty slots then  $\sum_{j=1}^{\infty} q_j = \sum_{j=1}^{\infty} j \cdot p_j$

## OPEN ADDRESSING – UNSUCCESSFUL FIND

### ○ Proof: (contd.)

- The expected number of probes is

$$1 + \sum_{j=0}^{\infty} j \cdot p_j = 1 + \sum_{j=1}^{\infty} q_j$$

- With uniform hashing

$$q_j = (n/m) * ((n-1)/(m-1)) * \dots * ((n-j+1)/(m-j+1)) \\ \leq (n/m)^j$$

- Then the expected number of probes is

$$1 + \sum_{j=1}^{\infty} q_j \leq 1 + \alpha + \alpha^2 + \alpha^3 + \dots \\ = 1 / (1 - \alpha)$$

## OPEN ADDRESSING – ANALYSIS - SUCCESSFUL FIND

- **Given:** open-address table with load factor  $\alpha = n/m < 1$
- **Assumptions:** Uniform Hashing; All keys are equally likely to be searched
- Expected number of probes in a successful find is at most  $1/\alpha + (1/\alpha) \cdot \ln(1/(1-\alpha))$
- **Proof:**
  - Simplifying assumption :
    - A successful find follows the same probe sequence as when the element was inserted
    - When is the assumption reasonable?
  - If  $k$  was the  $(j+1)^{\text{st}}$  key to be inserted
    - then the expected number of probes in finding  $k$  is given by the previous theorem (on unsuccessful find)
$$1/(1 - (j/m)) = m/(m-j)$$

# OPEN ADDRESSING – ANALYSIS - SUCCESSFUL FIND

## ○ Proof: (contd.)

- Expected number of probes in finding the key that was inserted as the  $(j+1)^{\text{st}}$  is  $m/(m-j)$
- Average over all  $n$  keys in the table

$$\begin{aligned} (1/n) \sum_{j=0}^{n-1} (m/(m-j)) &= (m/n) * (\sum_{j=0}^{n-1} (1/(m-j))) \\ &= (1/\alpha) * (H_m - H_{m-n}) \end{aligned}$$

where  $H_m$  is the  $m^{\text{th}}$  Harmonic number.

- Since  $\ln(j) \leq H_j \leq 1 + \ln(j)$

$$\begin{aligned} (1/\alpha) * (H_m - H_{m-n}) &\leq (1/\alpha) * (1 + \ln m - \ln(m-n)) \\ &= (1/\alpha) + (1/\alpha) * \ln(m/m-n) \\ &= 1/\alpha + (1/\alpha) * \ln(1/(1-\alpha)) \end{aligned}$$



# RE-HASHING

- Hash tables support efficient find operations:
  - Average case time complexity is  $O(1)$  if load factor is low
    - Load factor must be  $< 1$  for separate chaining
- In practice,
  - Load factor must be  $< 0.75$  to expect good performance.
- What if the hash table is nearly “full”?
  - Extend the hash table (i.e. increase its size)
    - Can the new hash function assign the old values to the same buckets as before?
      - bucket addresses must change for a good distribution?
  - Re-insert all the elements in the table
    - Referred to as *re-hashing*.

# RE-HASHING

- Cost of Rehashing
  - $O(\max(m,n))$  time – typically  $O(n)$  as table is nearly full.
    - Amortized Cost:  $O(1)$  time per element
  - But response time at the point of rehashing is bad:
    - allocation and copying of all the values takes  $O(n)$  time between two operations.
    - Or between the request for an operation and the response.
  - This is bad for applications requiring
    - bounded (worst case) response time
- What should be the size of the extended table?
  - Typical choice:  $2 * |T|$
  - Trade-offs: ???

# BLOOM FILTERS - MOTIVATION

- Tradeoff: Space vs. (In)Correctness
  - i.e. storage space for the table vs. false positives (membership)
- Example Problem: Stemming of words in search engine indexing:
  - e..g. plurals stemmed to singular; all parts of speech stemmed to one form
    - 90% of cases can be handled by simple rules
      - Rest - the exceptions – need a dictionary lookup
    - Suppose dictionary is large and must be stored in disk

## BLOOM FILTERS - MOTIVATION

- Consider this outline for stemming :

for each word  $w$

if ( $w$  is an exception word)

then  $\text{getStem}(w,D)$

else  $\text{apply-simple-rule}(w)$

} Need dictionary  
lookup on disk

- Cost for checking exceptions:

- $N * T_d$  where

- $N$  is # words and

- $T_d$  is lookup time (on disk)

# BLOOM FILTERS - MOTIVATION

- Suppose we can trade-off space for false positives (in lookup):  
for each word  $w$

```

if (w is in Dm )      // in-memory lookup (probabilistic)
then { s = getStem(w, Dd);  // disk lookup (deterministic)
    if invalid(s) then apply-simple-rule(w);
    } else { apply-simple-rule(w); }

```

- Cost for checking exceptions:

- $N * T_m + (r + f) * N * T_d$ 
  - $r$  is the proportion of exception words
  - $f$  is false positive rate
  - $T_m$  is lookup time in memory
  - $T_d$  is lookup time on disk
- Time Saved:  $(1 - r - f) * (T_d - T_m) / T_d$

## BLOOM FILTERS – AN IMPLEMENTATION

- Hash table is an array of bits indexed from 0 to  $m-1$ .
  - Initialize all bits to 0.
  - insert( $k$ ):
    - Compute  $h_1(k), h_2(k), \dots, h_d(k)$  where each  $h_i$  is a hash function resulting in one of the  $m$  addresses.
    - Set all those addressed locations to 1.
  - find( $k$ ):
    - Compute  $h_1(k), h_2(k), \dots, h_d(k)$
    - If all addressed locations are 1 then  $k$  is **found**  
Else  $k$  is **not found**
      - ↙  
**Always correct.**
      - ↘  
**Not necessarily correct!**

## BLOOM FILTERS - ANALYSIS

- Consider a table H of size m.
- Assume we use d “good” hash functions.
- After n elements have been inserted, the probability that *a specific location is 0* is given by
  - $p = (1 - 1/m)^{dn} \approx e^{-dn/m}$
- Let q be the proportion of 0 bits after insertion of n elements
  - Then the expected value  $E(q) = p$
- Claim (w/o proof):
  - With high probability q is close to its mean.
- So, the false positive rate is:
  - $f = (1-q)^d = (1-p)^d = (1 - e^{-dn/m})^d$

# BLOOM FILTERS

- The data structure is probabilistic:
  - If a value is not found then it is definitely not a member
  - If a value is found then it may or may not be a member.
- The error probability can be traded for space.
  - In practice, one can get low error probability with a (small) constant number of bits per element: (1 in our example implementation) .
- Applications:
  - Dictionaries (for spell-checkers, passwords, etc.)
  - Distributed Databases – exchange Bloom Filters instead of full lists.
  - Network Processing – Caches – exchange Bloom Filters instead of cache contents
  - Distributed Systems – P2P hash tables : instead of keeping track of all objects in other nodes, keep a Bloom filter for each node.



# LAS VEGAS VS. MONTE CARLO

- Quicksort:
  - Randomization for improved performance – correctness not altered
- Hashtables (for unordered dictionaries) :
  - Any 1-to-1 mapping will yield a table but a good hash function should yield a “uniformly random” distribution
  - Universal hashing chooses hash function “randomly”
- Both of the above are optimizations:
  - Such techniques are referred to as Las Vegas techniques.
- Monte Carlo Technique
  - Bloom Filter - Randomization yields a probabilistic algorithm that does not always produce correct results.

## DICTIONARY DATA STRUCTURES – SEARCH TREES

### Comparison of Sorted Arrays and Hashtables Ordered Dictionaries

- Better Representation
- Binary Trees
- Binary Search Trees
  - Implementation (Find, Add, and Delete)
  - Efficiency
  - Order Queries

### Balancing a Search Tree

- Height Balance Property

# DICTIONARY IMPLEMENTATIONS - COMPARISON

## Sorted Array

- Suitable for:
  - Ordered Dictionary
    - Example Queries: 2<sup>nd</sup> largest element? OR the element closest to k?
  - Offline operations (insertions/deletions)
  - Comparable Keys
- Implementation:
  - Deterministic

## Hashtable

- Suitable for:
  - Unordered Dictionary
  - Online insertions (deletions??)
    - Resizing can be done at an amortized cost of  $O(1)$  per element
  - Hashable Keys
- Implementation:
  - Randomized

# DICTIONARY IMPLEMENTATIONS - COMPARISON

## Sorted Array

- Time Complexity (find):
  - $\Theta(\log N)$  - worst case and average case
- Space Complexity
  - $\Theta(1)$

## Hashtable

- Time Complexity (find):
  - $\Theta(1)$  average case and  $\Theta(N)$  worst case
- Space Complexity
  - $\Theta(N)$  words – separate chaining (links)
  - $\Theta(N)$  bits – open addressing (empty & deleted flags)

## ORDERED DICTIONARY – BETTER REPRESENTATION?

- Is there an representation that
  - supports “relative order” queries and
  - supports online operations and
  - is resizable ?
- Revisit (the general structure of) Quicksort(Ls)

```
Quicksort(Ls) {
```

```
  If ( $|Ls| > 0$ ) {
```

```
    Partition Ls based on a pivot into LL and LG
```

```
    QuickSort LL
```


```
    QuickSort LG
```

```
  }
```

```
}
```

# ORDERED DICTIONARY – BETTER REPRESENTATION?

QuickSort  
Visualized



## ORDERED DICTIONARY – BETTER REPRESENTATION?

- Can we re-materialize the *QuickSort order* while searching?
  - i.e. a representation where key is compared with the pivot (pre-selected)
    - $\text{key} == \text{pivot} \implies \text{done}$
    - $\text{key} < \text{pivot} \implies \text{search in left subset}$
    - $\text{key} > \text{pivot} \implies \text{search in right subset.}$
- This is similar to QuickSelect but
  - With pre-selected pivots and stored “ordering” between the pivots.
    - i.e. ordering is preserved after sorting so as to support to “relative order” queries

# ORDERED DICTIONARY – BETTER REPRESENTATION?

- Data Model:

- A Set is characterized by the “Relation between Pivot and two (sub)sets”

- Generalized Data Model:

- A set is characterized by a “root” element and two subsets.



# ORDERED DICTIONARY – BETTER REPRESENTATION?

## ○ Inductive Definition:

- A binary tree is
  1. empty OR
  2. made of a root element and two binary trees - referred to as left and right (sub) trees
- For induction to be well founded “sub trees” must be of smaller size than the original.
- Sub trees are referred to as children (of the node which is referred to as the parent)
- A binary tree with two empty children is referred to as a leaf.

## ○ Inductive Definitions can be captured recursively:

BinaryTree = EmptyTree U (Element x BinaryTree x BinaryTree)

## ADT BINARY TREE

- BinaryTree createBinTree() // create empty tree
- Element getRoot(BinaryTree bt)
- BinaryTree getLeft(BinaryTree bt)
- BinaryTree getRight(BinaryTree bt)
- BinaryTree compose(Element root,  
BinaryTree leftBt,  
BinaryTree rightBt)

## ADT BINARY TREE - REPRESENTATION

- `struct __binTree;`
- `typedef struct __binTree *BinaryTree;`
- `struct __binTree {   Element rootVal;  
                          BinaryTree left;  
                          BinaryTree right;  
                          };`

Argue that the above representation in C captures the definition:

$\text{BinaryTree} = \text{EmptyTree} \cup (\text{Element} \times \text{BinaryTree} \times \text{BinaryTree})$

## ADT BINARY TREE - IMPLEMENTATION

```
BinaryTree compose(Element e, BinaryTree lt, BinaryTree rt)
{
    BinaryTree newT =
        (BinaryTree)malloc(sizeof(struct __binTree));
    newT->rootVal = e;
    newT->left = lt;
    newT->right = rt;
    return newT;
}
```

## ORDERED DICTIONARY – SEARCH TREE

- A binary search tree is
- a binary tree that captures an “ordering” (i.e. a relation)  $\mathcal{S}$  via the relation between the root and its subtrees:
  - i.e. for each element  $\underline{eL}$  in the left subtree:
    - $\underline{eL} \mathcal{S} \underline{rootVal}$
  - and for each element  $\underline{eR}$  in the right subtree:
    - $\underline{rootVal} \mathcal{S} \underline{eR}$

## ADT ORDERED DICTIONARY

- Element find(OrdDict d, Key k)
- OrdDict insert(OrdDict d, Element e)
- OrdDict delete(OrdDict d, Key k)
  - Note on Representation:
    - We can use the same BinaryTree representation for this.
      - i.e. The ordering is captured implicitly at the point of insertion by leveraging the left and right information.
    - Hence the following type definition – in C – would serve as the data definition!
  - End of Note.
- typedef BinaryTree OrdDict;

## ADT ORDERED DICTIONARY – IMPLEMENTATION

//Preconditions: k is unique;

Element find(OrdDict d, Key k)

```
{  
    if (d==NULL) return NOT_FOUND;  
    if (d->rootVal.key == k) return d->rootVal;  
    else if (d->rootVal.key < k) return find(d->right, k);  
    else /* d->rootVal.key > k */ return find(d->left, k);  
}
```

(Trivial) Exercise: **Modify implementation for multiple elements with the same key value.**

End of Exercise.

# ADT ORDERED DICTIONARY - IMPLEMENTATION

```
//Preconditions: d is non-empty; keys are unique (i.e. duplicates);
OrdDict insert(OrdDict d, Element e)
{
    if (d->rootVal.key < e.key) {
        if (d->right == NULL) { d->right = makeSingleNode(e); }
        else { insert(d->right, e); }
    } else {
        if (d->left == NULL) { d->left = makeSingleNode(e); }
        else { insert(d->left, e); }
    }
    return d;
} /* Exercise: Modify the top-level procedure to handle the case of the
    "empty tree".
    Modify the procedure to handle duplicates.
    End of Exercise. */
```



# ADT ORDERED DICTIONARY - IMPLEMENTATION

```
void makeSingleNode(Element e)
```

```
{
```

```
    OrdDict node;
```

```
    node = (OrdDict) malloc(sizeof(struct __binTree));
```

```
    node->rootVal=e;
```

```
    node->left = node->right = NULL;
```

```
    return node;
```

```
}
```

Exercise: Modify implementation for multiple elements with the same key (use one of the options):

- return success but do nothing,
- return failure with message “already found”,
- return success after adding new element separately,
- return success after overwriting contents.

End of Exercise

# ADT ORDERED DICTIONARY - IMPLEMENTATION

OrdDict delete(OrdDict dct, Key k)

- find the node, say nd, with contents matching key k
- if no such node exists done

else if nd is a leaf then delete nd // must free nd

else if one of the children of nd is empty

then replace nd with the other subtree of nd

else

in-order successor of nd will : (i) be within the  
subtree and (ii) have an empty left subtree

- a. find in-order successor of nd, say suc
- b. swap contents of suc with nd
- c. if suc is a leaf-node then delete suc // must free suc  
else replace suc with its right sub-tree

# ADT ORDERED DICTIONARY - IMPLEMENTATION

```
OrdDict delete(OrdDict dct, Key k)
{
    if (dct==NULL) return NULL;
    for (par=NULL, nd=dct; nd!=NULL; ) {
        if (nd->rootVal.key==k) break;
        else if (nd->rootVal.key < k) { par=nd; nd=nd->right;}
        else { par=nd; nd=nd->left; }
    }
    if (nd==NULL) return dct;
    if (par==NULL) { free(nd); return NULL; }
    else { return deleteSub(par, nd); }
}
```

# ADT ORDERED DICTIONARY - IMPLEMENTATION

```
OrdDict deleteSub(OrdDict par, OrdDict toDel) {  
    if (toDel->left!=NULL && toDel->right!=NULL) {  
        return deleteSubReplace(par, toDel);  
    } else if (toDel->right!=NULL) {  
        if (par->left==toDel) { par->left=toDel->right; }  
        else { par->right=toDel->right; }  
    } else if (toDel->left!=NULL) {  
        if (par->left==toDel) { par->left=toDel->left; }  
        else { par->right=toDel->left; }  
    } else {  
        if (par->left==toDel) {par->left=NULL;}  
        else {par->right=NULL;}  
    }  
    free(toDel); return dct;  
}
```

3/15/2014

find in-order successor of nd,  
say suc

- swap contents of suc with nd
- if suc is a leaf-node then delete suc // must free suc  
else replace suc with its right subtree

## ADT ORDERED DICTIONARY - IMPLEMENTATION

```
OrdDict deleteSubReplace(OrdDict par, OrdDict del)
{
    for (par=del,suc=del->right; suc->left!=NULL; par=suc,suc=suc-
        >left) ;
    swapContents(del, suc);
    if (suc->right==NULL) {
        if (par->left==suc) {par->left=NULL;}
        else {par->right=NULL; }
    } else {
        if (par->left==suc) { par->left=suc->right; }
        else { par->right=suc->right; }
    }
    free(suc); return dct;
}
```

# ADT ORDERED DICTIONARY - COMPLEXITY

- Time Complexity:
  - Find, insert, delete
    - Height of the tree
- Height of binary tree (by induction):
  - Empty Tree  $\Rightarrow 0$
  - Non-empty  $\Rightarrow 1 + \max(\text{height}(\text{left}), \text{height}(\text{right}))$
- Balanced Tree
  - Height =  $\log N$ 
    - Why?
- Unbalanced Tree
  - Worst case height =  $N$ 
    - Example?

# BINARY SEARCH TREES (BSTs)

- BSTs store data in order:
  - i.e. if you traverse a BST such that for all nodes  $v$ ,
    - Visit all nodes in the left sub tree of  $v$
    - Visit  $v$
    - Visit all nodes in the right sub tree of  $v$
  - then you are visiting them in sorted order.
- This is referred to as **in-order traversal**:

```
inorder(BinaryTree bt) {  
    if (bt != NULL) {  
        inorder(bt->left);  
        visit(bt);  
        inorder(bt->right);  
    }  
}  
// Time Complexity?? Space Complexity??
```

# BINARY SEARCH TREES (BSTs)

- Revisiting *delete* (in an Ordered Dictionary):
  - Deletion of an element with two non-empty subtrees required a pull-up operation.
  - One way of pulling-up –
    - find an element, say *c*, closest to the element to be deleted, say *d*
      - How?
    - overwrite *d* with *c*
    - delete node (originally) containing *c*
      - Will this result in recursive pulling-up? Why or why not?



## BINARY SEARCH TREES (BSTs)

### ○ Revisiting *delete* (in an Ordered Dictionary):

- Here is the ***pullUpLeft*** procedure

```
pullUpLeft(OrdDict toDel, OrdDict cur) {  
    pre = toDel;  
    while (cur->right != NULL) { pre=cur; cur=cur->right; }  
    toDel->rootVal = cur->rootVal;  
    if (cur->left==NULL) { prev->right = NULL; }  
    else { prev->right = cur->left; }  
    free(cur);  
}
```

// Exercise: Write a pullUpRight procedure

# BINARY SEARCH TREES – ORDER QUERIES

## ○ Exercises:

- Write a procedure to find the minimum element in a BST.
- Write a procedure to find the maximum element in a BST
- Write a procedure to find the second smallest element in a BST.
- Write a procedure to find the  $k^{\text{th}}$  smallest element in a BST.
- Write a procedure to find the element closest to a given element in a BST.

## ○ Hint:

- In all the above cases, use in-order traversal and stop once you get the result.

# BINARY SEARCH TREE - COMPLEXITY

- Time Complexity:
  - Find, insert, delete
    - # steps = Height of the tree
- Height of binary tree (by induction):
  - Empty Tree  $\Rightarrow 0$
  - Non-empty  $\Rightarrow 1 + \max(\text{height}(\text{left}), \text{height}(\text{right}))$
- Balanced Tree – Best case
  - Height =  $\log(N)$  where  $N$  is the number of nodes
- Unbalanced Tree – Worst case
  - Worst case height =  $N$  where  $N$  is the number of nodes
- How do you ensure balance?

## HEIGHT-BALANCE PROPERTY

- A node  $v$  in a binary tree is said to be *height-balanced* if
  - the difference between the heights of the children of  $v$  – its sub-trees – is at most 1.
- Height Balance Property:
  - A binary tree is said to be *height-balanced* if each of its nodes is height-balanced.
- Adel'son-Vel'skii and Landis tree (or AVL tree)
  - Any height-balanced binary tree is referred to as an AVL tree.
- The height-balance property keeps the height minimal
  - How?

# AVL TREE - HEIGHT

## ○ Theorem:

- The minimum number of nodes  $n(h)$  of an AVL tree of height  $h$  is  $\Omega(c^h)$  for some constant  $c > 1$ .

## ○ Proof (By induction):

1.  $n(1) = 1$  and  $n(2) = 2$
2. For  $h > 2$ ,  $n(h) \geq n(h-1) + n(h-2) + 1$   
Why?
3. Then,  $n(h)$  is a monotonic sequence i.e.  $n(h) > n(h-1)$ . So,  
 $n(h) > 2 * n(h-2)$
4. By, repeated substitution,  $n(h) > 2^j * n(h-2*j)$  for  $h-2*j \geq 1$
5. So,  $n(h)$  is  $\Omega(2^h)$

# AVL TREE - HEIGHT

- Corollary:
  - The height of an AVL tree with  $n$  nodes is  $O(\log n)$ .
  - Proof:
    - Obvious from the previous theorem.
- Thus the cost of a *find* operation in an AVL tree with  $n$  nodes is  $O(\log n)$ .
- What about insertion and deletion?
  - Adding or removing a node may disturb the balance.